

# A Static-Analysis Gate for Adaptive Retrieval: Cutting Structural Hallucinations in Repository-Level Code Completion

first last name  
Delft University of Technology  
Delft, Netherlands  
email@tudelft.nl

Second Author  
Delft University of Technology  
Delft, Netherlands  
email@tudelft.nl

**Abstract**—Retrieval-augmented generation (RAG) improves repository-level code completion, but retrieving on every request is wasteful and can even degrade output. Adaptive methods such as CARD gate retrieval on the model’s token-level confidence; this signal, however, cannot tell a syntactically fluent line from one that references identifiers that do not exist in the repository. We add a lightweight, training-free static-analysis gate on top of CARD’s confidence gate: when CARD decides to skip retrieval, we parse the zero-shot completion and trigger retrieval if it uses an out-of-scope (unresolvable) identifier. We evaluate the resulting cascade on four Python benchmarks spanning single-line (CrossCodeEval), multi-line function (RepoEval, CrossCodeLongEval), and fixed-block (CrossCodeLongEval) completion — 13,120 instances — with three code LMs (Qwen2.5-Coder-0.5B/1.5B and CodeLlama-7B), sweeping the retrieval budget end to end. Across all twelve (benchmark  $\times$  model) settings the static gate cuts the rate of invented identifiers (names absent from the ground truth that also resolve nowhere) by 3.6–11 $\times$  over CARD at a matched, single-digit-percent retrieval budget, while never lowering edit similarity (an asymmetry guarantee that holds at every threshold). We further find that always-retrieving does not reduce invented identifiers — it sometimes increases them — so the structural signal addresses a failure mode that neither raw retrieval nor confidence gating handles. We release code, the per-instance results, and a reproducible evaluation harness.

**Index Terms**—code completion, retrieval-augmented generation, adaptive retrieval, code language models, static analysis, hallucination, repository-level context

## I. Introduction

Repository level completion requires assistance from beyond the current file, as utility functions, shared variables and imported types can be defined elsewhere in the project. Retrieval Augmented Generation [1] addresses this file by augmenting the model’s prompt with cross-file-relevant snippets retrieved from the repository before generation. However, retrieval doesn’t always help due to pipelines either retrieving on every instance unconditionally and do not taking into account the local context or adaptive approaches relying too much on model’s confidence. Therefore, current approaches such as CARD [2] share a common issue: token-level confidence not distinguishing a syntactically correct line from one whose identifiers do not appear in the repository, not being

able to detect structural hallucinations. To address this gap, we augment current models’ confidence gate with a static analysis gate, which checks whether identifiers in significant positions are visible in the scope of the current file to be edited. The cascade pipeline operates in 3 stages: zero-shot generation, CARD’s confidence gate, and our free-training static analysis gate. The results of our experiments suggest that token-level confidence alone is insufficient and inefficient and that lightweight structural checks can help in avoiding structural hallucinations. Our main contribution consists of the static analysis gate approach which detects out-of-scope identifiers hallucinations without the need for fine tuning the model and the pipeline’s evaluation on the same settings as other RAG approaches such as CARD [2], enabling a reproducible evaluation and comparison.

## II. Problem Statement and Motivation

**Instruction:** Choose a project from the list provided to you. Clearly define the chosen task and explain why it intrigues your team (state why should people care about this topic).

### A. Research Questions

**Instruction:** Here it should be a story of RQs and how they relate and why we define them, rather than listing them. furthermore, we will have subsections for evaluation metrics, datasets (for both training testing), baselines, and configuration details. note that if dataset creation is an important step of your work, you can move it to the approach section.

Our study is organised around four questions that build on one another. We first ask whether retrieval is even worth gating: RQ1 — Is cross-file retrieval uniformly beneficial, or selective? If retrieval helped every instance, an always-retrieve policy would suffice and adaptivity would be pointless; if it is selective (or harmful on a subset), a gate is warranted. Given a gate is justified, we examine the state-of-the-art confidence gate: RQ2 — How well does CARD’s token-confidence gate trade retrieval budget for accuracy, and what failure mode does it leave?

We hypothesise that confidence correlates with accuracy but is blind to structural validity, letting confident yet unresolvable completions through. This motivates our contribution: RQ3 — Does a training-free static-analysis gate suppress the invented-identifier hallucinations that the confidence gate misses, and at what additional retrieval cost? Finally, because a gate that improves one metric by harming another is not useful, RQ4 — Does the static gate preserve accuracy and latency, and how do the effects scale with model size and task granularity? RQ1–RQ2 establish the setting and baseline, RQ3 measures the contribution, and RQ4 checks that it comes for free. The remainder of this section defines the evaluation metrics, datasets, baselines, and configuration used to answer them.

### III. Background and Related Work

**Instruction:** To ensure a comprehensive understanding and position your work in the broader landscape context, it is essential to familiarize yourself with recent advancements in the domain. Search for recent work that addresses your specific problem. Moreover, research studies that use your proposed model/method and summarise them. Use tools like Google Scholar, Semantic Scholar, or DBLP for your search. Prioritize contributions from top software engineering conferences such as ICSE, ESEC/FSE, ASE, and MSR and well-known journals like IEEE TSE, ACM TOSEM, and Springer EMSE. Find patterns in the relevant literature for instance work on memorization, LLMs, and attacks. etc. keep it short. make sure to highlight the gaps that you fill, what you add to the existing body of knowledge

#### A. Large Language Models (solution)

**Instruction:** Describe background knowledge, e.g., what is an LLM, fine-tuning, prompt engineering, chain of thought, instruct tuning, etc (any of which you use in your work)

#### B. Bug Localization (problem)

**Instruction:** Give relevant work on the problem you are solving, for example, you are proposing a bug localization approach, or you are working on benchmarks, etc. Study the relevant work and present them here.

#### C. Add more sub-headings based on your work

### IV. Approach

#### A. Design Rationale

CARD’s adaptive-retrieval gate [2] predicts the zero-shot output’s edit similarity from a 13-dimensional aggregation of per-token probabilities and entropies, and skips retrieval whenever this prediction exceeds a threshold  $T_{\text{RAG}}$ . The signal is cheap and empirically strong, but it is structurally shallow: token-level confidence cannot distinguish a syntactically correct line from one whose identifiers do not exist anywhere in the repository. We complement CARD with a second gate that parses the

structure of the zero-shot prediction rather than the model’s confidence, and triggers retrieval whenever the prediction falls under one of three classes of hallucination that the logit signal missed.

Because their decisions are not symmetric, we compose the two gates by subordination: the static gate runs only when CARD says skip. A static fire is positive evidence of a flagrant structural hallucination, while the absence of a static fire is not positive evidence of correctness, only that the prediction violates none of the patterns the gate enumerates. The static gate is therefore well-suited to supplement CARD’s skip decisions but cannot meaningfully veto its retrieve decisions. It is also naturally targeted at the residual failure mode CARD’s broad signal lets through: confidently-generated outputs that nonetheless contain a structural hallucination.

#### B. Pipeline

Given the in-file context  $X = (X_{\text{left}}, X_{\text{right}})$  surrounding the completion hole in the repository  $R$ , the cascade produces a prediction  $\hat{y}$  in three stages.

Stage 1: zero-shot generation. The code language model  $G$  produces  $\hat{y}_0 = G(X_{\text{left}}, X_{\text{right}})$  together with per-token log-probabilities. We extract CARD’s 13-D feature vector from the logits and compute  $\hat{s}_0 = \text{Estimator}(\cdot)$ , as proposed in the original paper.

Stage 2: CARD gate. If  $\hat{s}_0 < T_{\text{RAG}}$ , the model is judged uncertain. We retrieve cross-file context with BM25 over sliding-window chunks (20-line windows, stride 10, top- $k$ ) following the RepoCoder convention [3], regenerate, and return the retrieved-context prediction.

Stage 3: static-analysis gate. If CARD said skip, we parse  $\hat{y}_0$  with tree-sitter and run three independent checks against the repository’s symbol table. If any check fires, we retrieve and regenerate exactly as in Stage 2. Otherwise we commit  $\hat{y}_0$  as the final prediction.

#### C. Data

The primary evaluation surface is the Python split of CrossCodeEval [4], which contains roughly 2,700 line-completion instances drawn from approximately 420 GitHub repositories. The benchmark authors selected its source repositories to limit training-data overlap with the code language models available at the time of its release. Each instance is one JSONL record providing the in-file context ( $X_{\text{left}}$  from the prompt field,  $X_{\text{right}}$  from right\_context), the held-out completion (the groundtruth field), metadata identifying the upstream repository and target file, and a cross-file context block that is the only view of  $R$  available to either the BM25 retriever or our static-analysis gate.

We use the rg1\_bm25 variant of the dataset because it ships this block directly. The cross-file chunk represents a set of 5 pre-retrieved passages selected by BM25 from the instance-specific  $R$ . Each chunk is a contiguous 10-line window of a source file accompanied by its repo-relative

path. The benchmark authors pre-tile every file in the upstream project into non-overlapping windows of that size and, for each instance, score those windows against the last 10 lines of  $X_{\text{left}}$ , returning the top matches. This same chunk list serves both the retriever, when retrieval fires, and as the substrate for the repository symbol table consulted by the static gate. The symbol table is constructed afresh for each instance, pooling this instance’s chunks with those shipped for any other instance of the same repository. BM25 is not re-run at inference.

As a secondary surface we evaluate on **the function split of RepoEval [3], which contains 455 multi-line function-body completion instances** drawn from the 8 Python repositories. (We focus on function-level completion rather than the line/API splits because multi-line generation is where structural hallucinations have room to appear; see Section IV-G.) Rather than shipping a pre-retrieved chunk list per instance, as CrossCodeEval does, RepoEval ships the upstream repositories on disk in full. Thus, BM25 is run at inference time with 20-line sliding windows at stride 10 computed live from the source files, with the top ten matches returned per query. The symbol table consulted by our static-analysis gate is built afresh for each instance, immediately after the runner loads every Python file in the target repository into the per-instance file map, and therefore covers the full repository rather than a BM25-selected slice. We adopt these settings unchanged so our reimplementations of the CARD baseline remains directly comparable to its published numbers [2].

To stress-test the models on longer, more hallucination-prone targets, we additionally evaluate on CrossCodeLongEval [5], the long-form benchmark introduced with Repoformer, in both its function (multi-line body) and chunk (fixed 1–6 line block) Python tasks (5,000 instances each). CrossCodeLongEval ships, per instance, a pre-built left/right in-file context and a retrieved cross-file block, which we use exactly as for CrossCodeEval. All four evaluation surfaces and their statistics are presented together in Section IV-G.

#### D. Tier 1: Out-of-scope identifiers

Tier 1 targets the most common structural hallucination: a call or reference to a name that is not defined anywhere visible at the hole. We parse  $\hat{y}_0$  and collect identifiers appearing in significant positions — a call target, an attribute receiver, a subscript value, a class base, a bare decorator, an exception type, or a raise target — positions where the identifier must denote a real, bound object for the code to run. An identifier in such a position fires Tier 1 if it is not resolvable in the lexical scope of the current file (it is not a builtin, not imported, and not bound by a preceding definition, assignment, parameter, comprehension, with/except alias, etc., in  $X_{\text{left}}/X_{\text{right}}$ ). Tier 1 is purely a scope check and needs no repository symbol table. In our evaluation we realise this check

TABLE I

Evaluation surfaces. “GT lines” is the ground-truth completion length in lines; “single” is the fraction of single-line targets. Scoring mode is the span the prediction is truncated to before computing accuracy (Section V-B).

| Benchmark         | Task     | #Inst. | #Repos | GT lines |     |     | Scoring    |
|-------------------|----------|--------|--------|----------|-----|-----|------------|
|                   |          |        |        | med      | p99 | max |            |
| CrossCodeEval [4] | line     | 2665   | 471    | 1        | 1   | 1   | first line |
| RepoEval [3]      | function | 455    | 8      | 8        | 27  | 29  | body       |
| CCLong [5]        | function | 5000   | 944    | 9        | 30  | 31  | body       |
| CCLong [5]        | chunk    | 5000   | 1460   | 1        | 4   | 6   | line-count |

with pyflakes’ undefined-name analysis (F821) over the reconstructed file, which gives a binding-complete, false-positive-resistant version of the same signal; this is the gate isolated in all reported cascade results.

#### E. Tier 2: Signature mismatches

Tier 2 catches calls to functions that do exist in the repository but are invoked incorrectly. For every call in  $\hat{y}_0$  whose callee can be resolved in the repository symbol table, we compare the call site against the callee’s declared signature (arity and keyword names) and fire on an incompatibility (e.g., too many positional arguments, or an unknown keyword). Unlike Tier 1, Tier 2 consults the symbol table built from the cross-file context.

#### F. Tier 3: Wrong-origin imports

Tier 3 validates the provenance of imported names and module attributes against the repository’s module map: it fires when a symbol is imported from a module that does not export it, or an attribute is accessed on a module that does not define it. Tiers 2 and 3 are implemented in our in-house tree-sitter analyzer and broaden the gate’s coverage; because Tier 1 is by far the dominant and most robust signal, the experiments in this paper isolate the Tier 1 gate and leave a per-tier ablation to future work (Section VII-B).

#### G. Datasets

Our approach is training-free (the static gate parses code; the CARD estimator is the only learned component and is calibrated once per model, not per dataset), so our data requirement is entirely on the evaluation side: we need repository-level completion benchmarks that (i) ship the cross-file context our gate and the BM25 retriever consume, (ii) were curated to limit overlap with code LM training data, and (iii) span a range of completion granularities, since our central claim is about a failure mode (structural hallucination) whose prevalence we expect to grow with the length of the generated span. We therefore evaluate on four Python surfaces drawn from three established benchmarks; their statistics are summarised in Table I and Fig. 1.

Single-line: CrossCodeEval (line). Every target is exactly one line (median 38 characters); the model emits a short, mostly local completion. We include it as the easy




Fig. 1. Completion-target length by benchmark (log scale, box = IQR, whiskers 5–95th percentile). The four surfaces form a granularity spectrum: single-token/ single-line (CrossCodeEval), short fixed blocks (CCLong-chunk, 66% single-line), and multi-line function bodies (RepoEval, CCLong-function, median 8–9 lines).

end of the spectrum and as the setting used by our CARD baseline [2]: with so little generated code, there is little room for a structural hallucination, so this surface is a sanity floor on which we expect the static gate to fire rarely and change little.

Multi-line functions: RepoEval (function) and CCLong (function). Here the hole is an entire function body — a median of 8–9 lines and up to 31. This is the regime we find most interesting for two reasons. First, generating many lines gives the model many opportunities to reference helper functions, classes, or attributes that do not exist, so we expect structural hallucinations to be most frequent here — this is where we want to “push the models and see where they hallucinate.” Second, function completion is what lets us introduce truncation: a code LM asked for one function body keeps generating into the next definition, so the raw output must be cut at the end of the body before it is scored. We truncate at the dedent boundary of the generated body (the Python analogue of CARD’s bracket-matching), which materially changes the measured accuracy (Section VI) and, as we show, must be handled carefully so that it does not interact with the static check.

Fixed blocks: CCLong (chunk). The chunk task completes a fixed window of 1–6 lines (66% are single-line, median 1, max 6). It sits between the single-line and function regimes: short enough that targets are often a single statement, but multi-line often enough (34%) that the prediction must be truncated to the gold’s line count — matching Repoformer’s chunk metric — rather than to a function body. We use it to test whether our findings on functions degrade gracefully to shorter, less structurally rich targets.

Why this selection. The benchmarks were curated by their authors to reduce training-data contamination (CrossCodeEval samples permissively licensed repositories created after common pre-training cutoffs; CCLong is built from 1,500 held-out repositories), which protects the credibility of the no-retrieve baseline. More importantly

for our study, the four surfaces deliberately bracket the hypothesised effect: the smaller, shorter targets (CrossCodeEval line, CCLong chunk) are where we expect few hallucinations and therefore act as controls, while the longer function-body targets (RepoEval, CCLong function) are where we expect the static gate to matter most. Confirming that the gate is both near-inert on the controls and strongly beneficial on the long targets is exactly the evidence our claim needs.

## V. Experiment Setup

### A. Models and Configurations

We evaluate three open-weight code LMs spanning an order of magnitude in size: Qwen2.5-Coder-0.5B and Qwen2.5-Coder-1.5B [6] and CodeLlama-7B [7]. All are used zero-shot (no fine-tuning of the generator); the only learned component is CARD’s critique estimator. We compare four configurations, all sharing the same generator and retriever so differences isolate the gating policy: C1 no-retrieve (in-file context only), C2 always-retrieve, C3 CARD (retrieve iff  $\hat{s}_0 < T_{\text{RAG}}$ ), and C4 the cascade (C3, plus retrieve when the Tier-1 static gate fires on a skipped instance). Rather than reporting C3/C4 at one threshold, we sweep  $T_{\text{RAG}} \in \{0.05, \dots, 0.95\}$  to trace the entire cost–accuracy operating curve.

### B. Evaluation Setting and Metrics

Setting. All generation is greedy (temperature 0), so results are deterministic and re-runs are exact. Test/train overlap is mitigated upstream by the benchmarks’ curation (Section IV-G); we add no fine-tuning that could leak test data. Because a code LM asked for a multi-line body keeps generating past it, the prediction is truncated before accuracy is scored: to the first line for CrossCodeEval, to the generated function body (dedent boundary) for the function tasks, and to the ground-truth line count for chunks (matching Repoformer). We report every accuracy metric both on this truncated span (the headline) and on the raw, untruncated output, so the effect of truncation is visible rather than hidden.

Accuracy metrics. Exact Match (EM) is the fraction of predictions equal to the ground truth after truncation. Edit Similarity (ES) is  $1 - \text{Lev}(\hat{y}, y) / \max(|\hat{y}|, |y|) \in [0, 1]$ , the standard character-level metric used by CARD and CrossCodeEval [2], [4]. Identifier-F1 (idF1) is the set-F1 over identifier tokens, the identifier-matching metric of CrossCodeEval [4].

Hallucination metrics. Our target failure mode is the invented identifier: a name the model emits that is both (A4) absent from the ground truth and (B2) unresolvable by static analysis. We report  $h_{A4B2}$ , the fraction of predictions containing at least one identifier satisfying both conditions (B2 via pyflakes’ F821 undefined-name analysis on the reconstructed file), as the strict “made-up name” signal, and the looser  $h_{A4}$  (absent-from-gold only) as an upper bound. Crucially, the hallucination

metrics are always computed on the model’s complete (untruncated) output: truncating to an oracle length can delete a definition the rest of the file uses — thereby manufacturing an undefined-name flag that the model never produced — and the gold length is unavailable at inference. Computing the structural signal on the complete generation both avoids this artefact and matches what a deployed gate would see. The same principle governs the cascade’s Tier-1 trigger.

Efficiency metrics. We report the retrieval rate (% of instances for which retrieval fired), the direct cost knob, and per-instance latency (ms). For the adaptive configs latency is the cost model of CARD/Repoformer: a zero-shot pass for every instance plus a second retrieval-augmented pass only when the gate fires.

### C. Configuration and Implementation Details

Generation. We use fill-in-the-middle (FIM) prompting [8] with each model’s native FIM sentinels, plus a set of stop strings that halt generation at the natural boundary (preventing the `<|fim_pad|>` degeneration that otherwise floods short completions and corrupts both the text and CARD’s per-token features). The generation budget is set per task to the 99th percentile of gold length in tokens — 50 (CrossCodeEval), 280 (RepoEval-function), 400 (CCLong-function), and 80 (CCLong-chunk) — so the model can express the full target without unbounded over-generation.

Adaptive gate. CARD’s estimator is a LightGBM regressor [9] over a 13-dimensional aggregation of per-token probabilities and entropies that predicts the zero-shot output’s ES, reproduced from the original paper [2]. It is calibrated once per model on samples from The Stack (deduplicated) [10] and reused across all datasets (the estimator is a property of the generator, not the benchmark). The Tier-1 static gate is realised with pyflakes over the reconstructed file. Retrieval is Okapi BM25 [11] over 20-line/stride-10 windows, returning the top-10 chunks, following the RepoCoder convention [3].

Infrastructure and scale. All inference runs on a single NVIDIA A100-80GB GPU (Google Colab) using vLLM [12] with PagedAttention and continuous batching; the chip is fixed across models so the reported latencies are comparable within a benchmark. We batch the C1/C2 generation passes and share a generation cache across configurations; the full  $T_{\text{RAG}}$  sweep is then computed post hoc by replaying CARD and the cascade over the frozen generations, so no instance is generated twice. In total the study covers 13,120 instances  $\times$  two generation passes  $\times$  three models, plus the 19-point threshold sweep.

## VI. Results

We answer the four research questions in turn. Table II gives the two fixed-policy endpoints (C1 never retrieve, C2 always retrieve) for every benchmark and model; Table III isolates the cascade’s contribution against CARD at a low

retrieval budget; Table IV reports the truncated-vs-raw accuracy. Figures 3–6 visualise the trade-offs over the full threshold sweep.

### A. RQ1: Retrieval is broadly beneficial, but not uniformly

The C1→C2 columns of Table II show that adding cross-file context lifts edit similarity substantially on every surface and model — by +0.21 to +0.46 ES, with exact match rising several-fold (e.g. CCLong-function 1.5B: EM 0.08  $\rightarrow$  0.57). The no-retrieve baselines are thus “correctly low”: the models are knowledge-limited, not broken, which we confirmed by manual inspection (the C1 outputs are fluent but reference the wrong, locally plausible names). Figure 2 decomposes this per instance: on the function task, always-retrieve helps 77–80% of instances, leaves  $\sim$ 13% unchanged, and hurts 6–11%. So in our small-model, cross-file-dependent setting retrieval is more broadly useful than the  $\sim$ 20%-helpful regime reported for larger models by Repoformer [5] — but a non-trivial slice is still wasted or harmed, which is precisely the slice an adaptive policy should reclaim. Retrieval is therefore worth gating, motivating RQ2–RQ4.

### B. RQ2: CARD trades cost for accuracy well, but leaves a structural gap

Figure 3 traces accuracy against retrieval budget as  $T_{\text{RAG}}$  is swept. The cascade (and CARD, which nearly coincides on ES) recovers most of the C2 accuracy well before reaching a 100% retrieval rate: the curves rise steeply and then flatten, so a confidence gate buys most of retrieval’s benefit at a fraction of its cost — reproducing CARD’s and Repoformer’s central efficiency finding [2], [5] on our benchmarks and models. Figure 4 shows the same trade-off in wall-clock terms: edit similarity and per-instance latency both rise with  $T_{\text{RAG}}$  and plateau together. However, the confidence gate is blind to one thing: the  $h_{A4B2}$  columns of Table II show that always-retrieving does not reduce invented identifiers — on the function tasks at 7B it increases them (RepoEval nearly triples, 0.018  $\rightarrow$  0.051; CCLong 0.042  $\rightarrow$  0.054), because retrieval lets the model write more, longer code and thus emit more unresolvable names. A policy that only optimises confidence/ES therefore leaves a structural failure mode untouched. This is the gap our static gate targets.

### C. RQ3: The static gate suppresses the hallucinations CARD misses

Table III compares CARD (C3) and the cascade (C4) at a low budget ( $T_{\text{RAG}} = 0.05$ , where CARD is most conservative and retrieves almost nothing). Adding the Tier-1 static gate retrieves an extra 2–9% of instances — precisely those whose confident zero-shot output references an unresolvable name — and cuts the invented-identifier rate by 3.6 $\times$  to 11.3 $\times$  (to zero on RepoEval-1.5B). The effect is consistent across all twelve (benchmark  $\times$  model) cells. Figure 5 shows it holds along the entire budget



TABLE II

Endpoints: C1 (never retrieve) vs. C2 (always retrieve), headline (truncated) scoring. ES/EM/idF1 higher is better;  $h_{A4B2}$  (invented-identifier rate) lower is better; ms is per-instance latency under batched inference. Note that C2 lifts accuracy sharply everywhere, yet  $h_{A4B2}$  barely moves and often rises (e.g. RepoEval/CCLE-function at 7B): retrieval does not fix invented identifiers.

| Dataset         | M    | C1 no-retrieve |       |       |            |      | C2 always-retrieve |       |       |            |      |
|-----------------|------|----------------|-------|-------|------------|------|--------------------|-------|-------|------------|------|
|                 |      | ES             | EM    | idF1  | $h_{A4B2}$ | ms   | ES                 | EM    | idF1  | $h_{A4B2}$ | ms   |
| CCE (line)      | 0.5B | 0.556          | 0.205 | 0.518 | 0.0428     | 199  | 0.929              | 0.856 | 0.923 | 0.0398     | 264  |
|                 | 1.5B | 0.631          | 0.294 | 0.603 | 0.0210     | 200  | 0.961              | 0.918 | 0.958 | 0.0120     | 234  |
|                 | 7B   | 0.652          | 0.315 | 0.626 | 0.0169     | 502  | 0.961              | 0.923 | 0.958 | 0.0244     | 844  |
| RepoEval (func) | 0.5B | 0.357          | 0.046 | 0.483 | 0.0835     | 1280 | 0.717              | 0.339 | 0.784 | 0.0813     | 1248 |
|                 | 1.5B | 0.431          | 0.055 | 0.585 | 0.0527     | 1106 | 0.887              | 0.624 | 0.933 | 0.0505     | 1399 |
|                 | 7B   | 0.444          | 0.084 | 0.558 | 0.0176     | 3607 | 0.839              | 0.626 | 0.865 | 0.0505     | 4916 |
| CCLE (func)     | 0.5B | 0.397          | 0.039 | 0.520 | 0.0660     | 153  | 0.739              | 0.350 | 0.822 | 0.0754     | 186  |
|                 | 1.5B | 0.468          | 0.080 | 0.605 | 0.0342     | 170  | 0.849              | 0.574 | 0.906 | 0.0348     | 262  |
|                 | 7B   | 0.489          | 0.085 | 0.616 | 0.0418     | 904  | 0.839              | 0.600 | 0.890 | 0.0536     | 1594 |
| CCLE (chunk)    | 0.5B | 0.605          | 0.265 | 0.580 | 0.0208     | 58   | 0.846              | 0.649 | 0.849 | 0.0224     | 96   |
|                 | 1.5B | 0.707          | 0.394 | 0.690 | 0.0164     | 102  | 0.935              | 0.863 | 0.937 | 0.0148     | 174  |
|                 | 7B   | 0.711          | 0.419 | 0.691 | 0.0206     | 422  | 0.924              | 0.846 | 0.919 | 0.0240     | 725  |




Fig. 2. Per-instance effect of always-retrieve (1.5B):  $\Delta\text{ES} = \text{ES}(\text{C2}) - \text{ES}(\text{C1})$ . Retrieval is broadly beneficial on the function/chunk tasks, but a non-trivial fraction of instances is unchanged or degraded — the budget should be spent selectively rather than uniformly (RQ1).

axis: at any matched retrieval rate the cascade (solid) sits below CARD (dashed) in invented-identifier rate, the two converging only as both approach always-retrieve (where there is nothing left to select). Figure 6 summarises the headline reduction as paired bars. Because the trigger and the  $h_{A4B2}$  metric both rest on undefined-name analysis, part of this drop is structural by construction; we therefore also verified that the looser, fully independent  $h_{A4}$  (absent-from-gold) moves in the same direction, and that edit similarity does not fall (next subsection).

D. RQ4: The gate is free and scales as hypothesised

No accuracy regression. The cascade obeys an asymmetry guarantee by construction — it only adds retrieval to CARD’s skip decisions — and the data confirm it empirically: at every one of the 19 thresholds, in every cell, C4 retrieval rate  $\geq$  C3, C4 ES  $\geq$  C3, and C4  $h_{A4B2} \leq$  C3 (Table III shows ES ticking up, e.g. CCLong-function 0.5B 0.397  $\rightarrow$  0.424). The hallucination reduction is thus obtained at no cost to accuracy and only single-digit-percent extra retrieval (hence near-C1 latency).

Scaling. Accuracy rises monotonically with model size on every surface (Table II). Invented-identifier rates fall with size on the cleaner benchmarks (CCE C1  $h_{A4B2}$ : 0.043/0.021/0.017; RepoEval: 0.084/0.053/0.018 for 0.5B/1.5B/7B) but are noisier on the longer CCLong targets (7B-function 0.042  $\gtrsim$  1.5B 0.034) — longer, harder contexts make even the 7B invent names, which is exactly where the gate keeps paying off. The contribution is largest in absolute terms where hallucination is largest (smaller models, function tasks) and is near-inert on the single-line control (Fig. 6), as intended.

Truncation matters. Table IV shows the headline (truncated) ES is far above the raw ES on every multi-line surface (e.g. 7B CCLong-function C2 0.839 truncated vs. 0.426 raw): without body/line truncation, over-generation past the target span dominates the score. This both validates the truncation step and underlines why the structural signal must be computed on the complete output rather than the truncated one (Section V-B).

Accuracy-cost trade-off: the cascade reaches near-always-retrieve ES at a fraction of the retrieval budget




Fig. 3. Accuracy-cost trade-off (RQ2). Edit similarity vs. retrieval rate as  $T_{\text{RAG}}$  is swept (C4 cascade);  $\bullet$  = C1,  $\blacksquare$  = C2. The selective frontier reaches near-C2 accuracy well below a 100% retrieval budget, for all three models and all four surfaces.

TABLE III

RQ3 — the cascade’s contribution at a low budget ( $T_{\text{RAG}} = 0.05$ ). For a few percent of extra retrieval (r%), the Tier-1 static gate cuts the invented-identifier rate  $h_{A4B2}$  by the factor in the last column, while edit similarity is non-decreasing. “ $\rightarrow 0$ ” denotes a drop to zero.

| Dataset         | M    | C3 CARD |       |            |  | C4 cascade (ours) |       |            |                 |
|-----------------|------|---------|-------|------------|--|-------------------|-------|------------|-----------------|
|                 |      | r%      | ES    | $h_{A4B2}$ |  | r%                | ES    | $h_{A4B2}$ | $\downarrow$    |
| CCE (line)      | 0.5B | 0       | 0.558 | 0.0428     |  | 5                 | 0.580 | 0.0045     | 9.5 $\times$    |
|                 | 1.5B | 0       | 0.631 | 0.0210     |  | 2                 | 0.642 | 0.0019     | 11.1 $\times$   |
|                 | 7B   | 0       | 0.652 | 0.0169     |  | 2                 | 0.659 | 0.0015     | 11.3 $\times$   |
| RepoEval (func) | 0.5B | 0       | 0.357 | 0.0835     |  | 9                 | 0.391 | 0.0198     | 4.2 $\times$    |
|                 | 1.5B | 0       | 0.431 | 0.0527     |  | 6                 | 0.456 | 0.0000     | $\rightarrow 0$ |
|                 | 7B   | 0       | 0.444 | 0.0176     |  | 3                 | 0.460 | 0.0044     | 4.0 $\times$    |
| CCLE (func)     | 0.5B | 0       | 0.397 | 0.0660     |  | 8                 | 0.424 | 0.0134     | 4.9 $\times$    |
|                 | 1.5B | 0       | 0.468 | 0.0342     |  | 4                 | 0.489 | 0.0042     | 8.1 $\times$    |
|                 | 7B   | 0       | 0.489 | 0.0418     |  | 5                 | 0.504 | 0.0114     | 3.7 $\times$    |
| CCLE (chunk)    | 0.5B | 0       | 0.605 | 0.0208     |  | 3                 | 0.613 | 0.0038     | 5.5 $\times$    |
|                 | 1.5B | 0       | 0.707 | 0.0164     |  | 2                 | 0.713 | 0.0046     | 3.6 $\times$    |
|                 | 7B   | 0       | 0.711 | 0.0206     |  | 2                 | 0.718 | 0.0052     | 4.0 $\times$    |




Fig. 4. Accuracy and latency vs. the CARD threshold (7B cascade): both rise with  $T_{\text{RAG}}$  and plateau, so the threshold is a single knob trading per-instance latency for accuracy (RQ2).

## VII. Discussion

Confidence is not structure. The adaptive-retrieval line of work gates on the model’s own uncertainty: FLARE [13] re-retrieves when the next sentence contains low-probability tokens, and CARD [2] trains a critique model to predict the zero-shot output’s edit similarity from logit statistics. Our results make concrete a limitation

TABLE IV

RQ4 — truncated (headline) vs. raw edit similarity. Over-generation past the target span heavily deflates the raw score on every multi-line surface, justifying span truncation for accuracy (and on-complete scoring for hallucination).

| Dataset         | M    | C1 ES  |       | C2 ES  |       |
|-----------------|------|--------|-------|--------|-------|
|                 |      | trunc. | raw   | trunc. | raw   |
| RepoEval (func) | 0.5B | 0.357  | 0.265 | 0.717  | 0.475 |
|                 | 1.5B | 0.431  | 0.374 | 0.887  | 0.615 |
|                 | 7B   | 0.444  | 0.330 | 0.839  | 0.444 |
| CCLE (func)     | 0.5B | 0.397  | 0.287 | 0.739  | 0.472 |
|                 | 1.5B | 0.468  | 0.389 | 0.849  | 0.554 |
|                 | 7B   | 0.489  | 0.356 | 0.839  | 0.426 |
| CCLE (chunk)    | 0.5B | 0.605  | 0.478 | 0.846  | 0.564 |
|                 | 1.5B | 0.707  | 0.575 | 0.935  | 0.681 |
|                 | 7B   | 0.711  | 0.468 | 0.924  | 0.511 |

implicit in both: token-level confidence is necessary but not sufficient. A model can be perfectly confident in a call to a helper that does not exist — the tokens are individually likely because the shape is right — and exactly these




Fig. 5. Invented-identifier rate vs. retrieval budget (RQ3). At a matched budget the cascade (C4, solid) lies below CARD (C3, dashed) for every model; the gap is largest at the low budgets where a selective policy operates, and closes as both approach always-retrieve.

At a low retrieval budget ( $T_{\text{RAG}} = 0.05$ ), the static gate cuts invented-identifier hallucination several-fold over CARD (factor annotated), at a few % extra retrieval




Fig. 6. Headline reduction (RQ3): invented-identifier rate for CARD vs. the cascade at  $T_{\text{RAG}} = 0.05$ ; the annotated factor is the reduction. The gate is near-inert on the easy single-line control yet strongly beneficial on the function tasks, exactly as hypothesised.

cases slip through the confidence gate. The static check is cheap, training-free, and complementary: it reads the same zero-shot output CARD already produced and asks an orthogonal question (“does this resolve?”) rather than a confidence question (“is this likely?”).

Retrieval is not a hallucination cure. Repoformer [5] showed that always-retrieving is wasteful for accuracy — up to 80% of retrievals do not help. We add a sharper, complementary observation: always-retrieving does not fix structural hallucination and can worsen it (Table II), because more context invites more (and longer) generation, and thus more opportunities to invent names. This reframes selective retrieval: the goal is not only to skip unhelpful retrievals but to spend the budget on the structurally broken outputs, which is exactly what the Tier-1 gate does. It also explains why the cascade’s advantage is concentrated at low budgets (Fig. 5): once a policy retrieves almost everywhere, selectivity — and hence the gate’s targeting — has no room to act.

Why functions. The contribution tracks generation length: near-inert on single-line CrossCodeEval, strongest

on function bodies (Fig. 6). This is intuitive — a one-line completion rarely has space to call a non-existent function — and it is why benchmarks like CCLong [5], which deliberately lengthen the target, are the right testbed for structural hallucination. It also exposes a measurement subtlety the community should heed: a model asked for one function over-generates into the next, so accuracy needs span-truncation (Table IV) while the structural signal needs the complete output; conflating the two (truncating before the static check) can manufacture hallucinations that the model never produced.

#### A. Implications

For practitioners, the static gate is a near-zero-cost add-on to any confidence-gated RAG completion system: it needs no training, runs in milliseconds on the already-generated output, and removes the bulk of invented-identifier hallucinations at a few percent extra retrieval and no accuracy loss — a favourable trade in an IDE setting where a confidently-wrong API call is more damaging than a missing one. For the research community, our results argue that (i) selective-retrieval policies



should consider structural signals, not only confidence; (ii) repository-level completion should be evaluated with an explicit structural-hallucination metric, not edit similarity alone, since retrieval can improve ES while leaving (or worsening) hallucination; and (iii) such metrics must be computed on the model’s complete output, decoupled from the accuracy-side truncation. The full per-instance results and the harness that sweeps any gate over the budget axis are released to support this.

## B. Future Work

Several directions follow directly. First, a per-tier ablation: we isolated the Tier-1 (out-of-scope) gate; Tiers 2–3 (signature and import-origin checks, Sections IV-E–IV-F) should be measured separately, as they target rarer but higher-confidence structural errors. Second, repair instead of retrieve: the static analyzer already localises the offending identifier, so a cheaper action than full retrieval may be to constrain or correct that token. Third, better function-level critique labels: Repoformer notes that ES-based self-evaluation is miscalibrated for function completion, and our CARD estimator inherits this; a structural or execution-based label could sharpen the confidence gate so the cascade starts from a stronger base. Fourth, broader coverage: more model families and sizes, and languages beyond Python (the gate is language-agnostic in principle, needing only a parser and a scope resolver). Finally, repository-adaptive thresholds, since duplication and modularity vary across projects [3].

## C. Threats to the Validity

Construct. Edit similarity and exact match measure surface overlap, not functional correctness; we do not run tests (pass@k), so a “correct” completion may still be semantically wrong, and vice versa. Our hallucination metric is a static proxy: pyflakes cannot resolve dynamically-created or star-imported names, which could flag a valid identifier as undefined; we mitigate this by restricting Tier-1 to significant positions and by computing the signal on the reconstructed file (which contains the in-file imports and definitions). There is a partial circularity: the cascade’s trigger and the  $h_{A4B2}$  metric both rest on undefined-name analysis, so some of the measured reduction is structural by construction; we address this by also reporting the independent absent-from-gold rate  $h_{A4}$  (which moves the same way) and by showing edit similarity never decreases.

Internal. We re-implemented the CARD baseline rather than using original code, and reused the 1.5B/7B critique estimators across datasets; a calibration difference would shift C3 (and thus the C3→C4 gap) but not the asymmetry guarantee, which is structural. Generation is greedy and deterministic, so results are reproducible but reflect a single decode. Latencies are measured under vLLM continuous batching and are therefore throughput-amortised; we use them only for within-benchmark, same-hardware

comparison and rely on retrieval rate as the hardware-independent cost proxy. We report point estimates over large test sets (up to 5,000 instances); the harness supports paired McNemar and bootstrap tests, which we leave to a camera-ready version.

External. Findings cover three code LMs (0.5B–7B), Python only, and four benchmarks; larger models hallucinate less (Table II), so the absolute benefit may shrink at frontier scale even as the asymmetry guarantee persists. The CCLong-chunk surface strips blank lines during its construction, which does not affect Python scoping but makes its raw text differ slightly from on-disk source. Results on other languages or proprietary models may differ.

## VIII. Conclusion

We set out to close a structural blind spot in adaptive retrieval for repository-level code completion: confidence gates such as CARD skip retrieval on outputs that are fluent but reference identifiers that exist nowhere in the repository. Our remedy is a training-free static-analysis gate that runs only when the confidence gate decides to skip, parses the zero-shot completion, and triggers retrieval when it finds an out-of-scope identifier. Evaluated over four Python surfaces and three code LMs (13,120 instances), with the retrieval budget swept end to end, the cascade cuts the invented-identifier rate by 3.6–11× over CARD at a matched, single-digit-percent retrieval budget, never lowers edit similarity (an asymmetry guarantee that holds at every threshold), and addresses a failure mode that always-retrieving does not — and sometimes worsens. The effect is largest where it should be (smaller models, longer function targets) and near-inert on easy single-line completions. Beyond the specific gate, the study argues that selective retrieval should weigh structural evidence alongside confidence, and that repository-level completion should be measured with a hallucination metric computed on the model’s complete output. We release the code, the full per-instance results, and the sweeping evaluation harness to make these comparisons reproducible.

## References

- [1] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. tau Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive nlp tasks,” 2021. [Online]. Available: <https://arxiv.org/abs/2005.11401>
- [2] W. Zhang, T. Fu, T. Yuan, G. Zhang, D. Chen, and J. Wang, “A lightweight framework for adaptive retrieval in code completion with critique model,” 2024.
- [3] F. Zhang, B. Chen, Y. Zhang, J. Keung, J. Liu, D. Zan, Y. Mao, J.-G. Lou, and W. Chen, “RepoCoder: Repository-level code completion through iterative retrieval and generation,” in Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2023, pp. 2471–2484.
- [4] Y. Ding, Z. Wang, W. U. Ahmad, H. Ding, M. Tan, N. Jain, M. K. Ramanathan, R. Nallapati, P. Bhatia, D. Roth, and B. Xi-ang, “CrossCodeEval: A diverse and multilingual benchmark for cross-file code completion,” in Advances in Neural Information Processing Systems 36 (NeurIPS 2023) Track on Datasets and Benchmarks, 2023.

- [5] D. Wu, W. U. Ahmad, D. Zhang, M. K. Ramanathan, and X. Ma, “Repoformer: Selective retrieval for repository-level code completion,” in *Proceedings of the 41st International Conference on Machine Learning (ICML)*, ser. PMLR, vol. 235, 2024.
- [6] B. Hui, J. Yang, Z. Cui, J. Yang, D. Liu, L. Zhang, T. Liu, J. Zhang, B. Yu, K. Lu, K. Dang, Y. Fan, Y. Zhang, A. Yang, R. Men, F. Huang, B. Zheng, Y. Miao, S. Quan, Y. Feng, X. Ren, X. Ren, J. Zhou, and J. Lin, “Qwen2.5-Coder technical report,” 2024.
- [7] B. Rozière, J. Gehring, F. Gloeckle, S. Sootla, I. Gat, X. E. Tan, Y. Adi, J. Liu, R. Sauvestre, T. Remez, J. Rapin et al., “Code llama: Open foundation models for code,” 2023.
- [8] M. Bavarian, H. Jun, N. Tezak, J. Schulman, C. McLeavey, J. Tworek, and M. Chen, “Efficient training of language models to fill in the middle,” 2022.
- [9] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, “LightGBM: A highly efficient gradient boosting decision tree,” in *Advances in Neural Information Processing Systems* 30 (NeurIPS), 2017.
- [10] D. Kocetkov, R. Li, L. Ben Allal, J. Li, C. Mou, C. Muñoz Ferrandis, Y. Jernite, M. Mitchell, S. Hughes, T. Wolf, D. Bahdanau, L. von Werra, and H. de Vries, “The stack: 3 TB of permissively licensed source code,” 2022.
- [11] S. Robertson and H. Zaragoza, “The probabilistic relevance framework: BM25 and beyond,” *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–389, 2009.
- [12] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica, “Efficient memory management for large language model serving with PagedAttention,” in *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, 2023.
- [13] Z. Jiang, F. F. Xu, L. Gao, Z. Sun, Q. Liu, J. Dwivedi-Yu, Y. Yang, J. Callan, and G. Neubig, “Active retrieval augmented generation,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.

## IX. Appendix: Task Assignment Sheet

Instruction: After discussion among yourself within your group, list the main project-related tasks assigned to and done by each team member. Tasks can be: conducting a literature review, preparing a dataset, implementation of technique X, code review, fine-tuning a model, prompt design, setting up eval pipeline, evaluating the technique X, improving repo documentation, writing section X of the report, preparing the presentation (section methodology, section evaluation, etc.). The goal is to have a balanced task distribution among team members.